

# MP1 Design Document

## Protocol

Despite the overhead from SYN/ACK and FIN/ACK from TCP; I chose TCP instead of UDP as I did not wish to implement my own system to deal with packet loss and arrival order.

On the topic of TCP, I set the `TCP_NODELAY` and `TCP_CORK` socket options with `setsockopt()`.

With `TCP_NODELAY` our data gets sent immediately, which is good since our payload is not much larger than the headers. However, this left us susceptible to silly window syndrome and congestion collapse; these massively decreased the performance of the program in my tests.

After testing with `TCP_CORK`, I found a significant improvement in throughput over Nagle's and `TCP_NODELAY`; however, the latency from the 200ms ACK delay was very noticable.

## Client-Server Communication

**Command Mode** Each command from the client begins with a 32-bit `MessageType`:

```
enum MessageType { CREATE, DELETE, JOIN, LIST, RESPONSE };
```

The `CREATE`, `DELETE`, `JOIN`, `LIST` types will be sent from client to server. These commands will be followed by a variable length null-terminated string from the client capped at 224 bytes.

The `RESPONSE` type sent from server to client. The data that follows the `RESPONSE` value depends on the client's message type:

- `CREATE` and `DELETE` is followed by a single 32-bit value from the `Status` enum
- `JOIN` is followed by two 32-bit values: `port` and `members`
- `LIST` is followed a null-terminated string.

For `LIST`, it would be better to send an integer with the string length before the string so we can know exactly how many bytes to read, thus improving performance and reliability; but I ran out of time to implement this.

**Chat Mode** After sending the `JOIN` message, the client will await for the port number from the `RESPONSE` message from the server and establish a new connection on said port. Now in chat mode, the client will wait for user input; upon receiving input the client will send a variable length null-terminated string over the socket. The chat thread associated with this client will then multicast the message to the clients subscribed to the chatroom.

## Server

**Parallelization** The server on the main thread will have a socket binded to the port specified on the command line. After accepting a client connection, the server creates a new thread to handle command messages from the client.

The command thread upon receiving `JOIN` message will create a new chatroom thread accepting connections to the room. When a client connects to the chatroom, a chat thread is created to handle chat messages from the client.

After reading about C10k and profiling with callgrind, I realized that the server wasted a lot of time through context switching between the chat threads. I attempted to use a single thread to `accept()` connections and `recv()/send()` chat messages with `epoll_wait()`. Stress testing with 6 GiB/s of input caused the TCP buffer to saturate causing `EAGAIN`; , the application was unable to recover after this. Due to the assignment deadline I did not have time to explore this issue, and instead I reverted to the multithreaded approach.

**Database** In an attempt to improve performance, I used the `std::unordered_map` to get  $O(1)$  access. `unordered_map` actually incurs a performance loss in the provided test cases due to the relatively large constant involved with the hashing function.

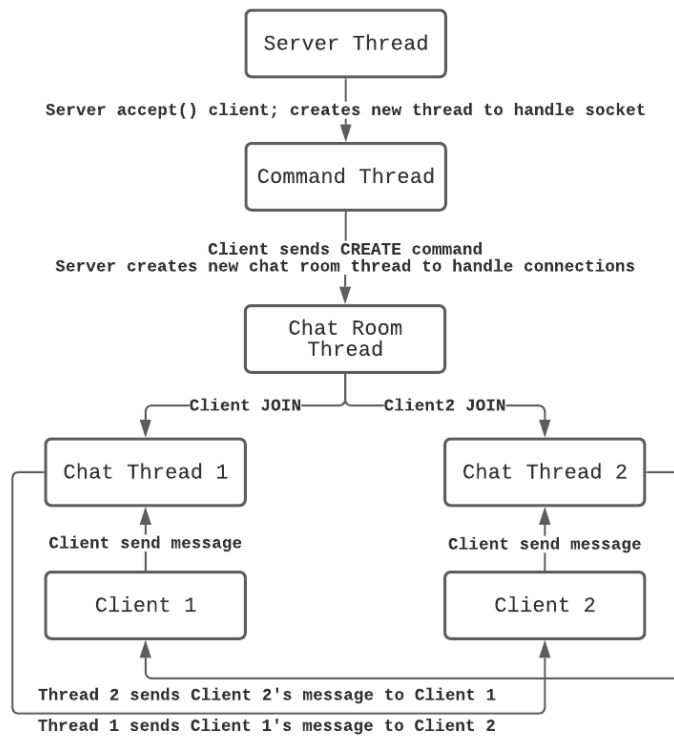


Figure 1: A diagram of the server behavior in response to commands from two connected clients.

## Client

**Chat Parallelization** The client's chat mode uses two threads: one for reading from the socket, and one for reading from `stdin`.

I found this approach to be faster than using `epoll_wait()` on a singular thread; even if I enabled edge triggers with `EPOLLET`.

I used `select()` in the user input loop so that I can check to see if the connection is alive, if not then I can kill the thread immediately. Without this, `fgets()` will block the thread from dying until the user enters a new line character.

## Known Issues

If we run `echo "create r1" | ./crc localhost 8080`, echo will close the pipe emitting an EOF, causing `get_command()` to run infinitely. This is because `get_message()` and `get_command()` in `interface.h` do not check if `fgets()` returns NULL.

I wanted to point this out as I don't know how the test script will work and I cannot modify `interface.h`.